

Attribute Information:

1. **name** - ASCII subject name and recording number
2. **mdvp_fo_hz** - Average vocal fundamental frequency (Actually column name MDVP:Fo(Hz))
3. **mdvp_fhi_hz** - Maximum vocal fundamental frequency (Actually column name MDVP:Fhi(Hz))
4. **mdvp_flo_hz** - Minimum vocal fundamental frequency (Actually column name MDVP:Flo(Hz))
5. **mdvp_jitter_in_percent, mdvp_jitter_abs, mdvp_rap, mdvp_ppq, jitter_ddd** - Several measures of variation in fundamental frequency (Actually column names MDVP:Jitter(%), MDVP:Jitter(Abs), MDVP:RAP, MDVP:PPQ, Jitter:DDP respectively)
6. **mdvp_shimmer, mdvp_shimmer_db, shimmer_apq3, shimmer_apq5, mdvp_apq, shimmer_ddd** - Several measures of variation in amplitude (Actually column names MDVP:Shimmer, MDVP:Shimmer(dB), Shimmer:APQ3, Shimmer:APQ5, MDVP:APQ, Shimmer:DDA respectively)
7. **nhr, hnr** - Two measures of ratio of noise to tonal components in the voice (Actually column names NHR, HNR respectively)
8. **rpde, d2** - Two nonlinear dynamical complexity measures (Actually column names RPDE, D2 respectively)
9. **dfa** - Signal fractal scaling exponent (Actually column name DFA)
10. **spread1, spread2, ppe** - Three nonlinear measures of fundamental frequency variation (Actually column names spread1, spread2, PPE respectively)
11. **status** - Health status of the subject (one) - Parkinson's, (zero) - healthy (**Target Variable / attribute**)

```
# Importing the necessary libraries
import numpy                as np
import pandas               as pd
import seaborn              as sns
import matplotlib.pyplot    as plt
import warnings

%matplotlib inline
warnings.filterwarnings('ignore')

import statsmodels.api      as sm
from sklearn import model_selection
from sklearn.model_selection import train_test_split

# getting methods for confusion matrix, F1 score, Accuracy Score
from sklearn import metrics
from sklearn.metrics        import confusion_matrix, f1_score, accuracy_score, classification_report, roc_curve, auc, average_precision_score
from sklearn.linear_model   import LogisticRegression
from sklearn.naive_bayes    import GaussianNB
from sklearn.neighbors      import KNeighborsClassifier
from sklearn.svm            import SVC

from sklearn.preprocessing import StandardScaler
```

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import StackingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
```

```
pdDataOrg = pd.read_csv("Parkinson disease.csv")
pdDataOrg.head()
```

	name	MDVP:F0(Hz)	MDVP:F1(Hz)	MDVP:F2(Hz)	MDVP:Jitter(%)	MDVP:Jitter(Abs)	MDVP:RAP
0	phon_R01_S01_1	119.992	157.302	74.997	0.00784	0.00007	0.00370
1	phon_R01_S01_2	122.400	148.650	113.819	0.00968	0.00008	0.00465
2	phon_R01_S01_3	116.682	131.111	111.555	0.01050	0.00009	0.00544
3	phon_R01_S01_4	116.676	137.871	111.366	0.00997	0.00009	0.00502
4	phon_R01_S01_5	116.014	141.781	110.655	0.01284	0.00011	0.00655

5 rows × 24 columns

```
pdData = pdDataOrg.copy()

targetCol = 'status'
targetColDf = pdData.pop(targetCol)
pdData.insert(len(pdData.columns),targetCol, targetColDf)

# deleting variables that were used for changing column position of target column
del targetCol
del targetColDf

# converting column names into lower case
pdData.columns = [c.lower() for c in pdData.columns]
# replacing spaces in column names with '_'
pdData.columns = [c.replace(' ', '_') for c in pdData.columns]
# replacing ':' in column names with '_'
pdData.columns = [c.replace(':', '_') for c in pdData.columns]
# replacing '(' in column names with '_'
pdData.columns = [c.replace('(', '_') for c in pdData.columns]
# replacing ')' in column names with '' i.e blank
pdData.columns = [c.replace(')', '') for c in pdData.columns]
# replacing '%' in column names with 'in_percent'
pdData.columns = [c.replace('%', 'in_percent') for c in pdData.columns]

# to check the above printing top 5 rows
pdData.head()
```

	name	mdvp_fo_hz	mdvp_fhi_hz	mdvp_flo_hz	mdvp_jitter_in_percent	mdvp_jitter_abs	mdvp_
0	phon_R01_S01_1	119.992	157.302	74.997	0.00784	0.00007	0.00
1	phon_R01_S01_2	122.400	148.650	113.819	0.00968	0.00008	0.00
2	phon_R01_S01_3	116.682	131.111	111.555	0.01050	0.00009	0.00
3	phon_R01_S01_4	116.676	137.871	111.366	0.00997	0.00009	0.00
4	phon_R01_S01_5	116.014	141.781	110.655	0.01284	0.00011	0.00

5 rows × 24 columns

```
print('\033[1mThe Parkinson\'s disease dataset having "{0}" rows and "{1}" columns\033[0m.'.format(pdData.shape[0],pdData.shape[1]))
```

The Parkinson's disease dataset having "195" rows and "24" columns.

```
pdData.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 195 entries, 0 to 194
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   name                   195 non-null    object
1   mdvp_fo_hz             195 non-null    float64
2   mdvp_fhi_hz            195 non-null    float64
3   mdvp_flo_hz            195 non-null    float64
4   mdvp_jitter_in_percent 195 non-null    float64
5   mdvp_jitter_abs        195 non-null    float64
6   mdvp_rap               195 non-null    float64
7   mdvp_ppq               195 non-null    float64
8   jitter_ddp             195 non-null    float64
9   mdvp_shimmer           195 non-null    float64
10  mdvp_shimmer_db        195 non-null    float64
11  shimmer_apq3           195 non-null    float64
12  shimmer_apq5           195 non-null    float64
13  mdvp_apq               195 non-null    float64
14  shimmer_dda            195 non-null    float64
15  nhr                    195 non-null    float64
16  hnr                    195 non-null    float64
17  rpde                   195 non-null    float64
18  dfa                    195 non-null    float64
19  spread1                195 non-null    float64
20  spread2                195 non-null    float64
21  d2                     195 non-null    float64
22  ppe                    195 non-null    float64
23  status                 195 non-null    int64
```

```
dtypes: float64(22), int64(1), object(1)
memory usage: 36.7+ KB
```

```
# setting name column as index column
pdData.set_index('name', inplace=True)
```

```
# after setting column 'name' as index now we have less columns to confirm that printing number of rows and column once again
print('\033[1mAfter setting \'name\' column as index of the Dataset,\033[0m now there are \033[1m"{0}"\033[0m Rows and \033[1m"{1}"\033[0m Columns in the given Dataset.'.format(pdData
```

After setting 'name' column as index of the Dataset, now there are "195" Rows and "23" Columns in the given Dataset.

```
pdData.head()
```

	mdvp_fo_hz	mdvp_fhi_hz	mdvp_flo_hz	mdvp_jitter_in_percent	mdvp_jitter_abs	mdvp_ra
name						
phon_R01_S01_1	119.992	157.302	74.997	0.00784	0.00007	0.0037
phon_R01_S01_2	122.400	148.650	113.819	0.00968	0.00008	0.0046
phon_R01_S01_3	116.682	131.111	111.555	0.01050	0.00009	0.0054
phon_R01_S01_4	116.676	137.871	111.366	0.00997	0.00009	0.0050
phon_R01_S01_5	116.014	141.781	110.655	0.01284	0.00011	0.0065

5 rows × 23 columns

```
# printing datatypes of each columns of the dataset
```

```
print("\033[1m***100)
print("a.\nColumn_Names      Data_Types")
print("***30)
print("\033[0m{0}\033[1m".format(pdData.dtypes))
print("***30)
print()
```

```
# printing No of Columns having different Types of Datatype
```

```
print("***100)
print("b.\nNumber of Columns with each DataTypes as follows :)")
print("***50)
print("Column_Names      No_of_Columns\033[0m")
print("***30)
print(pdData.dtypes.value_counts())
print("\033[1m***30)
```

```

print("\033[0m")

# printing Different Column Names of the dataset

print("\033[1m"*100)
print("c.\nEach Column Names of the dataset")
print(""*80)
print("\033[0m{0}\033[1m".format(pdData.columns))
print(""*80)
print("\033[0m")

*****

a.
Column_Names      Data_Types
*****
mdvp_fo_hz        float64
mdvp_fhi_hz       float64
mdvp_flo_hz       float64
mdvp_jitter_in_percent float64
mdvp_jitter_abs   float64
mdvp_rap          float64
mdvp_ppq          float64
jitter_ddp        float64
mdvp_shimmer      float64
mdvp_shimmer_db   float64
shimmer_apq3      float64
shimmer_apq5      float64
mdvp_apq          float64
shimmer_dda       float64
nhp               float64
hnp               float64
rpde              float64
dfa               float64
spread1           float64
spread2           float64
d2                float64
ppe               float64
status            int64
dtype: object
*****

b.
Number of Columns with each DataTypes as follows :
*****
Column_Names      No_of_Columns
*****
float64           22
int64              1
dtype: int64
*****

c.
*****

```

Each Column Names of the dataset

```
*****
Index(['mdvp_fo_hz', 'mdvp_fhi_hz', 'mdvp_flo_hz', 'mdvp_jitter_in_percent',
      'mdvp_jitter_abs', 'mdvp_rap', 'mdvp_ppq', 'jitter_ddp', 'mdvp_shimmer',
      'mdvp_shimmer_db', 'shimmer_apq3', 'shimmer_apq5', 'mdvp_apq',
      'shimmer_dda', 'nhf', 'hnr', 'rpde', 'dfa', 'spread1', 'spread2', 'd2',
      'ppe', 'status'],
      dtype='object')
*****
```

After observing the dataset and column description given we can conclude the followings:

- Columns having only two datatypes, int64, float64. (column 'name' was object datatype which was set as index of the dataframe)
- Column 'status' is only having int64 datatype, remaining all columns datatype is float64.
- All columns except 'status' are Numeric column.
- Columns 'status' is Nominal Categorical column with binary response.

```
# checking missing values in dataset for each attributes / columns
```

```
print("\033[1m"*100)
print("Column_Name      No_of_Missing_Values")
print("*****50)
print("\033[0m{0}".format(pdData.isnull().sum()))
print("\033[1m"*50)
print()
```

```
# checking if any duplicate rows available in the dataset
```

```
print("*****100)
print("Showing Duplicate rows if any in the dataset: ")
print("*****50)
print("\033[0m{0}".format(pdData[pdData.duplicated()])))
print("\033[1m"*100)
print("\033[0m")
```

```
*****
```

```
Column_Name      No_of_Missing_Values
```

```
*****
```

```
mdvp_fo_hz      0
mdvp_fhi_hz      0
mdvp_flo_hz      0
mdvp_jitter_in_percent  0
mdvp_jitter_abs  0
mdvp_rap         0
mdvp_ppq         0
jitter_ddp       0
mdvp_shimmer     0
mdvp_shimmer_db  0
shimmer_apq3     0
```

```

shimmer_apq5      0
mdvp_apq          0
shimmer_dda       0
nhr               0
hnr               0
rpde              0
dfa               0
spread1           0
spread2           0
d2                0
ppe               0
status            0
dtype: int64
*****

```

```

*****

```

Showing Duplicate rows if any in the dataset:

```

*****

```

Empty DataFrame

Columns: [mdvp_fo_hz, mdvp_fhi_hz, mdvp_flo_hz, mdvp_jitter_in_percent, mdvp_jitter_abs, mdvp_rap, mdvp_ppq, jitter_ddp, mdvp_shimmer, mdvp_shimmer_db, shimmer_apq3, shimmer_apq5, shimmer_dda, shimmer_hnr, shimmer_nhr, shimmer_rpd, shimmer_spr, shimmer_spr2, shimmer_d2, shimmer_ppe, shimmer_status]

[0 rows x 23 columns]

```

*****

```

As shown above,

(a.) There are no missing values

and (b.) No duplicate rows in the given dataset

```

# Five point summary of each attribute
pdData.describe().T

```

	count	mean	std	min	25%	50%	75%	
mdvp_fo_hz	195.0	154.228641	41.390065	88.333000	117.572000	148.790000	182.769000	260.10
mdvp_fhi_hz	195.0	197.104918	91.491548	102.145000	134.862500	175.829000	224.205500	592.03
mdvp_flo_hz	195.0	116.324631	43.521413	65.476000	84.291000	104.315000	140.018500	239.17
mdvp_jitter_in_percent	195.0	0.006220	0.004848	0.001680	0.003460	0.004940	0.007365	0.03
mdvp_jitter_abs	195.0	0.000044	0.000035	0.000007	0.000020	0.000030	0.000060	0.00
mdvp_rap	195.0	0.003306	0.002968	0.000680	0.001660	0.002500	0.003835	0.02
mdvp_ppq	195.0	0.003446	0.002759	0.000920	0.001860	0.002690	0.003955	0.01
jitter_ddp	195.0	0.009920	0.008903	0.002040	0.004985	0.007490	0.011505	0.06
mdvp_shimmer	195.0	0.029709	0.018857	0.009540	0.016505	0.022970	0.037885	0.11
mdvp_shimmer_db	195.0	0.282251	0.194877	0.085000	0.148500	0.221000	0.350000	1.30
shimmer_apq3	195.0	0.015664	0.010153	0.004550	0.008245	0.012790	0.020265	0.05
shimmer_apq5	195.0	0.017878	0.012024	0.005700	0.009580	0.013470	0.022380	0.07
mdvp_apq	195.0	0.024081	0.016947	0.007190	0.013080	0.018260	0.029400	0.13

```
# checking skewness of the data
pdData.skew().sort_values(ascending=False)
```

```

nhr          4.220709
jitter_ddp   3.362058
mdvp_rap     3.360708
mdvp_jitter_in_percent  3.084946
mdvp_ppq     3.073892
mdvp_jitter_abs  2.649071
mdvp_apq     2.618047
mdvp_fhi_hz  2.542146
mdvp_shimmer_db  1.999389
shimmer_apq5  1.798697
mdvp_shimmer  1.666480
shimmer_dda   1.580618
shimmer_apq3  1.580576
mdvp_flo_hz   1.217350
ppe          0.797491
mdvp_fo_hz   0.591737
spread1      0.432139
d2           0.430384
spread2      0.144430
dfa          -0.033214
rpde         -0.143402
hnr          -0.514317
status       -1.187727
dtype: float64
```


As from above we understand the following:

- Independent variables are measured in different units e.g. Hz, dB, % and absolute etc i.e variation in units of data exists and gap between feature values extremely high. Requires data scaling techniques to scale different quantities of measurements.
- Symmetrical distribution : Values close to 0 MDVP:F0(Hz) spread1 spread2 PPE
- Negative skewness and Tail is larger towards the left hand side of the distribution HNR status RPDE DFA
- Positive skewness and Tail is larger towards the Right hand side of the distribution All other attributes have a very high distribution towards right of the median

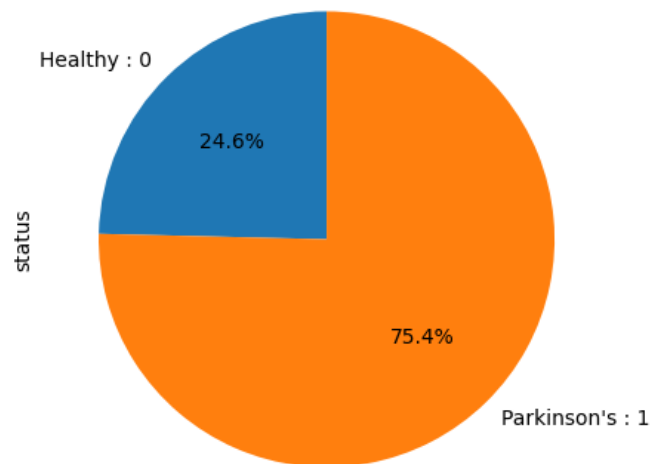
```
plt.figure(figsize=(10,5))                # setting figure size with width = 10 and height = 5
# seaborn count catplot to examine distribution of the status
ax = sns.catplot(x='status', kind="count", data=pdData)
plt.title("Distribution of column : 'Status'")    # setting title of the figure
y = []                                           # creating a null or empty array
for val in range(pdData.status.nunique()):      # looping for number of unique values in the status
    # appending count of each unique values from status to array y
    y.append(pdData.groupby(pdData.status,sort=False)['status'].count()[val])
for i, v in enumerate(y):                      # looping count of each unique value in the status
    # including count of each unique values in the plot
    plt.annotate(str(v), xy=(i,float(v)), xytext=(i-0.1, v+3), color='black', fontweight='bold')
```

<Figure size 1000x500 with 0 Axes>

```
plt.figure(figsize=(5,5))                # setting figure size with width = 10 and height = 5
# seaborn pie chart to examine distribution of the status
pdData.groupby(['status']).status.count().plot(kind='pie',labels=['Healthy : 0','Parkinson\'s : 1'],
                                              startangle=90, autopct='%1.1f%%')
plt.title("Distribution of column : 'status'") # setting title of the figure
```

```
Text(0.5, 1.0, "Distribution of column : 'status'")
```

Distribution of column : 'status'



From above we can see out of 195 patients, 48 patients (24.6 %) are healthy and 147 patients (75.4%) patients are having Parkinson's disease.

```
#Split the data into training and test set in the ratio of 70:30 respectively
X = pdData.drop(['status'],axis=1)
y = pdData['status']

# split data into train subset and test subset
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=47)

# checking the dimensions of the train & test subset
# printing dimension of train set
print(X_train.shape)
# printing dimension of test set
print(X_test.shape)
```

```
(136, 22)
(59, 22)
```

```
X_train.drop(['mdvp_jitter_in_percent'],axis=1,inplace=True)
X_test.drop(['mdvp_jitter_in_percent'],axis=1,inplace=True)
```

```
X_train.drop(['mdvp_shimmer'],axis=1,inplace=True)
X_test.drop(['mdvp_shimmer'],axis=1,inplace=True)
```

```
X_train.drop(['hnr'],axis=1,inplace=True)
X_test.drop(['hnr'],axis=1,inplace=True)
```

```
# printing dimension of train set
print(X_train.shape)
# printing dimension of test set
print(X_test.shape)
```

```
(136, 19)
(59, 19)
```

```
# Let us scale train as well as test data using StandardScaler
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.fit_transform(X_test)
```

```
# Train and Fit model
lr = LogisticRegression(random_state=0)
lr.fit(X_train_scaled, y_train)
```

```
#predict status for X_test_scaled dataset
lr_y_pred = lr.predict(X_test_scaled)
```

```
# Confusion Matrix for the Logistic Regression Model
print("Confusion Matrix : Logistic Regression")
print(confusion_matrix(y_test,lr_y_pred))
```

```
# Classification Report for the Logistic Regression Model
classRep = classification_report(y_test, lr_y_pred, digits=2)
print(classRep)
```

```
Confusion Matrix : Logistic Regression
[[ 9  4]
 [ 4 42]]
      precision    recall  f1-score   support

0         0.69      0.69      0.69         13
```

1	0.91	0.91	0.91	46
accuracy			0.86	59
macro avg	0.80	0.80	0.80	59
weighted avg	0.86	0.86	0.86	59

From the above Logistic Regression Model, we can find out the following details:

- **Accuracy of the model:- 86%**
- **Re-call of the model:- 91%**
- **Precision of the model:- 91%**
- **F1-Score of the model:- 91%**

```
# creating odd list of K for KNN
myList = list(range(3,40,2))

# creating empty list for F1 scores od different value of K
f1ScoreList = []

# perform accuracy metrics for values from 3,5....29
for k in myList:
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train_scaled, y_train)
    # predict the response
    y_pred = knn.predict(X_test_scaled)
    # evaluate F1 Score
    f1Score = f1_score(y_test, y_pred)
    f1ScoreList.append(f1Score)

# changing to misclassification error
MSE = [1 - x for x in f1ScoreList]

# determining best k
bestk = myList[MSE.index(min(MSE))]
print("The optimal number of neighbors is %d" % bestk)
```

The optimal number of neighbors is 29

```
# instantiate learning model (k = 29)
knn = KNeighborsClassifier(n_neighbors = 29, weights = 'uniform', metric='euclidean')

# fitting the model
knn.fit(X_train_scaled, y_train)

# predict the response
knn_y_pred = knn.predict(X_test_scaled)
```

```
# Confusion Matrix for the K-nearest neighbors Model
print("Confusion Matrix : K-nearest neighbors")
print(confusion_matrix(y_test,knn_y_pred))

# Classification Report for the K-nearest neighbors Model
classRep = classification_report(y_test, knn_y_pred, digits=2)
print(classRep)
```

```
Confusion Matrix : K-nearest neighbors
[[ 8  5]
 [ 0 46]]

              precision    recall  f1-score   support

     0         1.00      0.62      0.76         13
     1         0.90      1.00      0.95         46

 accuracy          0.92
 macro avg          0.95      0.81      0.86
 weighted avg       0.92      0.92      0.91
```

From the above K-nearest neighbors Model, we can find out the following details:

- **Accuracy of the model:- 92%**
- **Re-call of the model:- 100%**
- **Precision of the model:- 90%**
- **F1-Score of the model:- 95%**

```
svm = SVC(gamma=0.05, C=70,random_state=47)
svm.fit(X_train_scaled , y_train)

# predict the response
svm_y_pred = svm.predict(X_test_scaled)

# Confusion Matrix for the Support Vector Machine Model
print("Confusion Matrix : Support Vector Machine")
print(confusion_matrix(y_test,svm_y_pred))

# Classification Report for the Support Vector Machine Model
classRep = classification_report(y_test, svm_y_pred, digits=2)
print(classRep)
```

```
Confusion Matrix : Support Vector Machine
[[10  3]
 [ 0 46]]

              precision    recall  f1-score   support

     0         1.00      0.77      0.87         13
     1         0.94      1.00      0.97         46
```

accuracy			0.95	59
macro avg	0.97	0.88	0.92	59
weighted avg	0.95	0.95	0.95	59

From the above Support Vector Machine Model, we can find out the following details:

- **Accuracy of the model:- 95%**
- **Re-call of the model:- 100%**
- **Precision of the model:- 94%**
- **F1-Score of the model:- 97%**

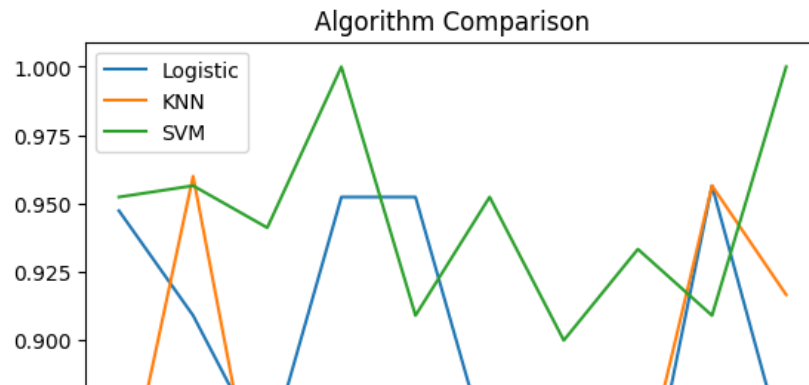
```
#Using K fold to check how the above algorithms varies throughout the dataset with 10 different subset of equal bins
models = []
models.append(('Logistic Regression', LogisticRegression(random_state=47)))
models.append(('K-NN', KNeighborsClassifier(n_neighbors = 29, weights = 'uniform', metric='euclidean')))
models.append(('SVM', SVC(gamma=0.05, C=70, random_state=47)))

# evaluate each model
results = []
names = []
scoring = 'f1'
for name, model in models:
    kfold = model_selection.KFold(n_splits=10, random_state=47, shuffle=True)
    cv_results = model_selection.cross_val_score(model, X_train_scaled, y_train, cv=kfold, scoring=scoring)
    results.append(cv_results)
    names.append(name)
    print("\033[1m{}\033[0m model have \033[1mmean F1-Score\033[0m of {1} and \033[1mSD F1-Score\033[0m of {2}".format(name, cv_results.mean(), cv_results.std()))
```

Logistic Regression model have **mean F1-Score** of 0.8930882775410977 and **SD F1-Score** of 0.05346316304648212
K-NN model have **mean F1-Score** of 0.8805069690035365 and **SD F1-Score** of 0.04633386023424363
SVM model have **mean F1-Score** of 0.9453975265995727 and **SD F1-Score** of 0.033233692576697435

```
plt.title('Algorithm Comparison')
plt.plot(results[0], label='Logistic')
plt.plot(results[1], label='KNN')
plt.plot(results[2], label='SVM')
plt.legend()
```

<matplotlib.legend.Legend at 0x7f2578d65430>



From the above comparison of different algorithms (Logistic Regression, K-nearest neighbors and Support Vector Machine) we can conclude that SVM (Support Vector Machine) performed slightly better than other algorithms.

```
# defining level heterogeneous model
level0 = list()
level0.append(('lr', LogisticRegression(random_state=47)))
level0.append(('knn', KNeighborsClassifier(n_neighbors = 29, weights = 'uniform', metric='euclidean')))
level0.append(('cart', DecisionTreeClassifier()))
level0.append(('svm', SVC(gamma=0.05, C=70, random_state=47)))
level0.append(('bayes', GaussianNB()))

# define meta learner model
level1 = SVC(gamma=0.05, C=3, random_state=47)

# define the stacking ensemble with cross validation of 5
Stack_model = StackingClassifier(estimators=level0, final_estimator=level1, cv=5)

# predict the response
Stack_model.fit(X_train_scaled, y_train)
prediction_stack = Stack_model.predict(X_test_scaled)

# Confusion Matrix for the Stacking Model
print("Confusion Matrix : Stacking")
print(confusion_matrix(y_test, prediction_stack))

# Classification Report for the Stacking Model
print(classification_report(y_test, prediction_stack, digits=2))
```

```
Confusion Matrix : Stacking
[[10  3]
 [ 0 46]]


      precision    recall  f1-score   support

0         1.00      0.77      0.87         13
```

	1	0.94	1.00	0.97	46
accuracy				0.95	59
macro avg		0.97	0.88	0.92	59
weighted avg		0.95	0.95	0.95	59

```
#determining false positive rate and True positive rate, threshold
fpr, tpr, threshold = metrics.roc_curve(y_test, prediction_stack)
roc_auc_stack = metrics.auc(fpr, tpr)
```

```
#plotting ROC curve
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc_stack)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



From the above Stacked meta classifier Model, we can find out the following details:

- Accuracy of the model:- 95%

- **Re-call of the model:- 100%**
- **Precision of the model:- 94%**
- **F1-Score of the model:- 97%**
- **ROC-AUC : 88%**

```
#creating model of Random Forest
RandomForest = RandomForestClassifier(n_estimators = 100,criterion='entropy',max_features=10,random_state=47)
RandomForest = RandomForest.fit(X_train_scaled, y_train)

# predict the response
RandomForest_pred = RandomForest.predict(X_test_scaled)

# Confusion Matrix for the Random Forest Model
print("Confusion Matrix : Random Forest")
print(confusion_matrix(y_test,RandomForest_pred))

# Classification Report for the Random Forest Model
print(classification_report(y_test, RandomForest_pred, digits=2))
```

```
Confusion Matrix : Random Forest
[[ 9  4]
 [ 1 45]]


              precision    recall  f1-score   support

     0           0.90       0.69       0.78         13
     1           0.92       0.98       0.95         46

 accuracy              0.92         59
 macro avg              0.91       0.84       0.86         59
 weighted avg           0.91       0.92       0.91         59
```

```
#determining false positive rate and True positive rate, threshold
fpr, tpr, threshold = metrics.roc_curve(y_test, RandomForest_pred)
roc_auc_rf = metrics.auc(fpr, tpr)
```

```
#plotting ROC curve
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc_rf)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```



From the above Random Forest Model, we can find out the following details:

- **Accuracy of the model:- 92%**
- **Re-call of the model:- 98%**
- **Precision of the model:- 92%**
- **F1-Score of the model:- 95%**
- **ROC-AUC : 84%**

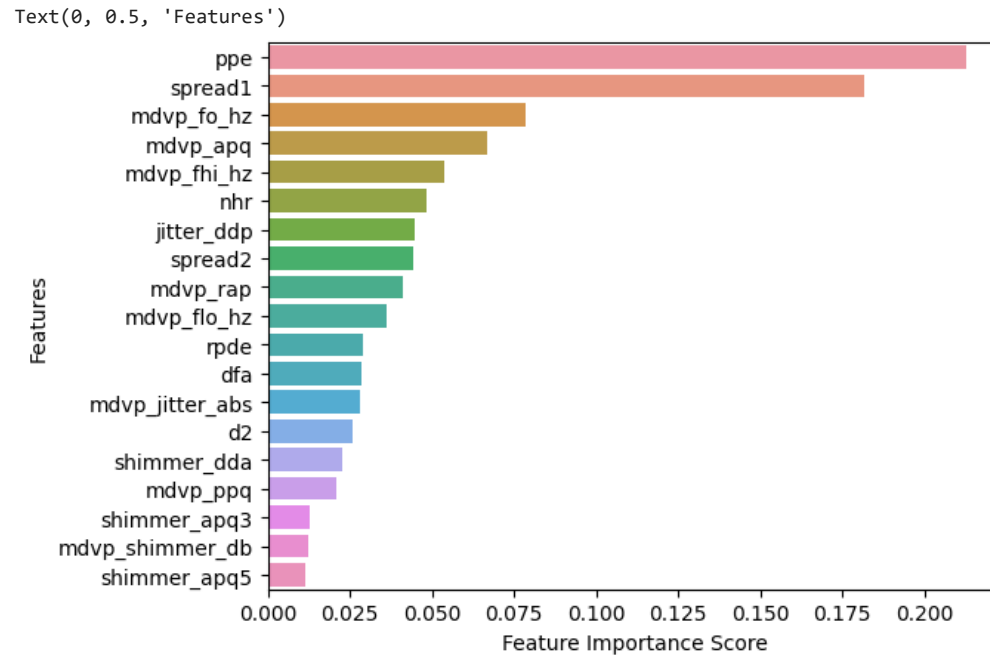
Lets check features importance

```
feature_imp = pd.Series(RandomForest.feature_importances_, index=X_train.columns).sort_values(ascending=False)
feature_imp
```

ppe	0.212646
spread1	0.181815
mdvp_fo_hz	0.078539
mdvp_apq	0.066849
mdvp_fhi_hz	0.053630
nhr	0.048251
jitter_ddp	0.044719
spread2	0.044320
mdvp_rap	0.041004
mdvp_flo_hz	0.036010
rpde	0.029126
dfa	0.028516
mdvp_jitter_abs	0.028306
d2	0.025957
shimmer_dda	0.022860

```
mdvp_ppq          0.020985
shimmer_apq3      0.012780
mdvp_shimmer_db   0.012384
shimmer_apq5      0.011302
dtype: float64
```

```
# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
```



```
#creating model of Adaptive Boosting
AdBs = AdaBoostClassifier( n_estimators= 50)
AdBs = AdBs.fit(X_train_scaled, y_train)

# predict the response
AdBs_y_pred = AdBs.predict(X_test_scaled)

# Confusion Matrix for the Adaptive Boosting Model
print("Confusion Matrix : Adaptive Boosting")
print(confusion_matrix(y_test,AdBs_y_pred))

# Classification Report for the Adaptive Boosting Model
print(classification_report(y_test, AdBs_y_pred, digits=2))
```

Confusion Matrix : Adaptive Boosting

```
[[ 9  4]
 [ 2 44]]
```

	precision	recall	f1-score	support
0	0.82	0.69	0.75	13
1	0.92	0.96	0.94	46
accuracy			0.90	59
macro avg	0.87	0.82	0.84	59
weighted avg	0.89	0.90	0.90	59

```
#determining false positive rate and True positive rate, threshold
fpr, tpr, threshold = metrics.roc_curve(y_test, AdBs_y_pred)
roc_auc_ada = metrics.auc(fpr, tpr)
```

```
#plotting ROC curve
plt.title('Receiver Operating Characteristic')
plt.plot(fpr, tpr, 'b', label = 'AUC = %0.2f' % roc_auc_ada)
plt.legend(loc = 'lower right')
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0, 1])
plt.ylim([0, 1])
plt.ylabel('True Positive Rate')
plt.xlabel('False Positive Rate')
plt.show()
```

Receiver Operating Characteristic

From the above Adaptive Boosting Model, we can find out the following details:

- **Accuracy of the model:- 90%**
- **Re-call of the model:- 96%**
- **Precision of the model:- 92%**
- **F1-Score of the model:- 94%**
- **ROC-AUC : 82%**

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=47)
```

```
# Train the model
from xgboost import XGBClassifier
```

```
model=XGBClassifier()
model.fit(X_train,y_train)
predict=model.predict(X_test)
```

```
print(accuracy_score(y_test,predict)*100)
```

```
# Confusion Matrix for the XGBoosting Model
print("Confusion Matrix : XGBoosting")
print(confusion_matrix(y_test,predict))
```

```
# Classification Report for the XGBoosting Model
print(classification_report(y_test, predict, digits=2))
```

```
97.43589743589743
```

```
Confusion Matrix : XGBoosting
```


```
[[ 7  1]
 [ 0 31]]
```

	precision	recall	f1-score	support
0	1.00	0.88	0.93	8
1	0.97	1.00	0.98	31
accuracy			0.97	39
macro avg	0.98	0.94	0.96	39
weighted avg	0.98	0.97	0.97	39

```
from sklearn.metrics import confusion_matrix
cm=confusion_matrix(y_test,predict)
plt.figure(figsize=(8,6))
fg=sns.heatmap(cm,annot=True,cmap="Reds")
figure=fg.get_figure()
```

```
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title("Output Confusion Matrix")
```

```
Text(0.5, 1.0, 'Output Confusion Matrix')
```



```
from sklearn.model_selection import KFold
from sklearn.model_selection import cross_val_score
# The baseline

# Test options and evaluation metric
num_folds = 10
seed = 7
scoring = 'accuracy'

# Spot-Check Algorithms
models = []
models.append(('Logistic Regression', LogisticRegression(random_state=47)))
models.append(('K-NN', KNeighborsClassifier(n_neighbors = 29, weights = 'uniform', metric='euclidean')))
models.append(('SVM', SVC(gamma=0.05, C=70, random_state=47)))
models.append(('Stacking', StackingClassifier(estimators=level0, final_estimator=level1, cv=5)))
```


```
models.append(('Random Forest', RandomForestClassifier(n_estimators = 100,criterion='entropy',max_features=10,random_state=47)))
models.append(('Adaptive Boosting', AdaBoostClassifier( n_estimators= 50)))
models.append(('XGBoost',XGBClassifier()))

def eval_algorithms(models, show_boxplots=True):
    # Evaluate each model in turn
    # Setup the test harness to use 10-fold cross validation
    results = []
    names = []
    for name, model in models:
        kfold = KFold(n_splits=10, random_state=seed,shuffle=True)
        cv_results = cross_val_score(model, X_train, y_train, cv=kfold, scoring=scoring)
        results.append(cv_results)
        names.append(name)
        #print("Estimated accuracy of {} with the mean of {} and std. dev. {}".format(name, cv_results.mean()*100.0, cv_results.std()*100.0))
        #print("{}: {} ({}).format(name, cv_results.mean()*100.0, cv_results.std()*100.0))

    if show_boxplots:
        # Create a plot of the model evaluation results to compae the spread
        # and the estimated mean accuracy of each model
        fig = plt.figure(figsize=(14,12))
        fig.suptitle('Algorithm Comparison')
        ax = fig.add_subplot(111)
        plt.boxplot(results)
        ax.set_xticklabels(names)
        plt.show()

eval_algorithms(models)
```

Algorithm Comparison

[Colab paid products](#) - [Cancel contracts here](#)